

**CSCI 4060U – Laboratory #8**  
**Programming Matrix Multiplication in OpenCL**  
Lab Due: 11pm, Apr. 11, 2023 (Canvas)

### Introduction

The main goal of this lab is to take a sequential implementation of matrix multiplication and parallelize it in OpenCL.

### Activity #1

First, we will start with a sequential implementation of matrix multiplication – [http://www.sqrlab.ca/exercises/csci4060u-w23/matrix\\_mult\\_seq.c](http://www.sqrlab.ca/exercises/csci4060u-w23/matrix_mult_seq.c). This program includes a function to multiple two matrices stored as one-dimensional arrays:

```
// [AB] = [A] * [B]
void seq_matrix_mult(
    int m_dim, int k_dim, int n_dim,
    float *a, float *b, float *ab);
```

The function implements matrix multiplication as follows:

$$\begin{aligned} \text{Matrix A} * \text{Matrix B} &= \text{Matrix AB} \\ [m * k] * [k * n] &= [m * n] \end{aligned}$$

Make sure you familiarize yourself with the sequential algorithm. Next, in OpenCL C write a parallel version of the algorithm to run on the device (GPUs) where every kernel calculates one element in Matrix AB. Finally, write the host code in C before testing your program. Compare the results of the sequential program with your OpenCL program to ensure that it works correctly.

### Activity #2

Although the parallel version is an improvement, it is not the most efficient parallelization. In a comment at the top of your code explain why and propose (no code needed!) an alternative parallelization scheme.

### Marking Scheme

#### Activity #1

- |                        |         |
|------------------------|---------|
| - Device (kernel) code | 3 marks |
| - Host code            | 3 marks |

#### Activity #2

- |   |         |
|---|---------|
| - Description of efficiency problem       | 2 marks |
| - Proposed improvement to parallelization | 2 marks |

<b>TOTAL</b>	<b>10 MARKS</b>
--------------	-----------------